Deep dive into Transformers

Amirhassan Amirmahani - Gholamreza Dar

Transformer Architecture

- Transformer architecture
- Encoder and Decoder stack
- Self-attention
- Multihead attention
- Positional encoding
- Residual connections
- Encoder - Decoder Attention
- Encoder Attention Vs Decoder Attention (Masking)
- The final Linear and Softmax layers
- Training
  - One Hot Encoding
  - Loss Function
  - Another Training Method

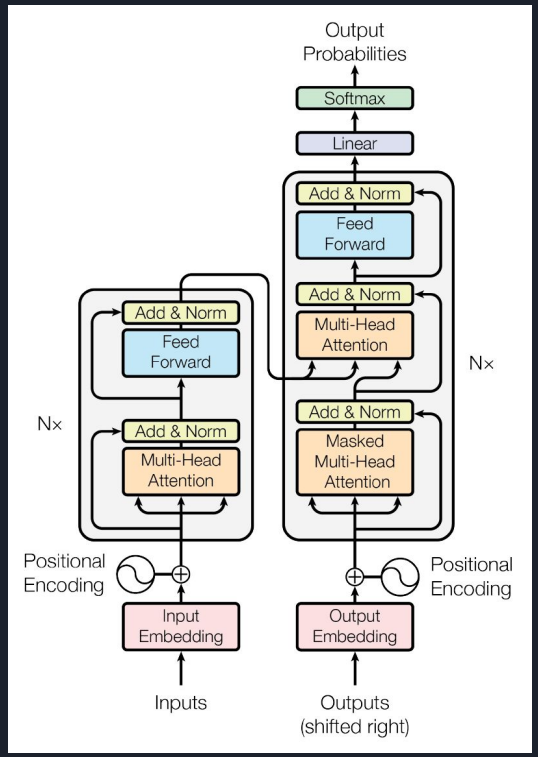

Transformer as a black box

http://jalammar.github.io/illustrated-transformer/

Inside a Transformer

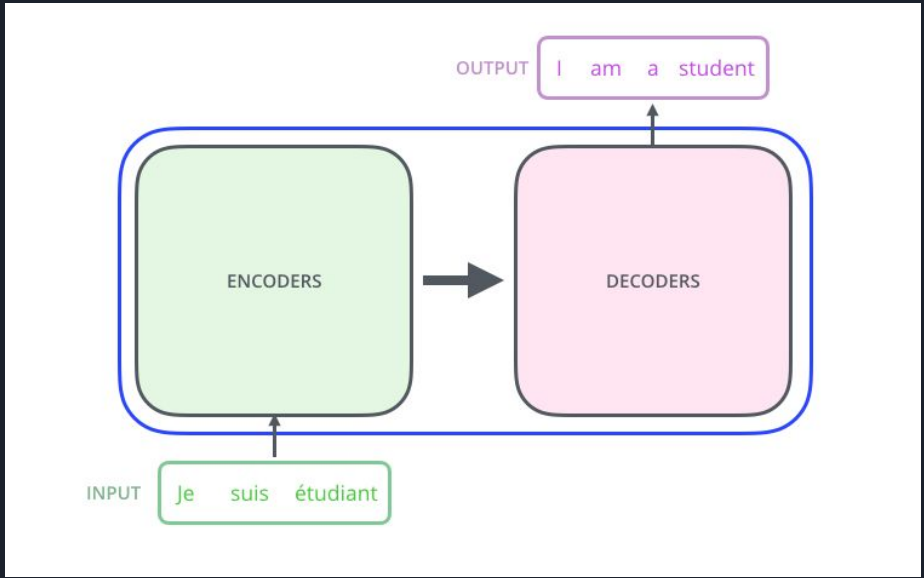

http://jalammar.github.io/illustrated-transformer/

Encoder and Decoder Stacks

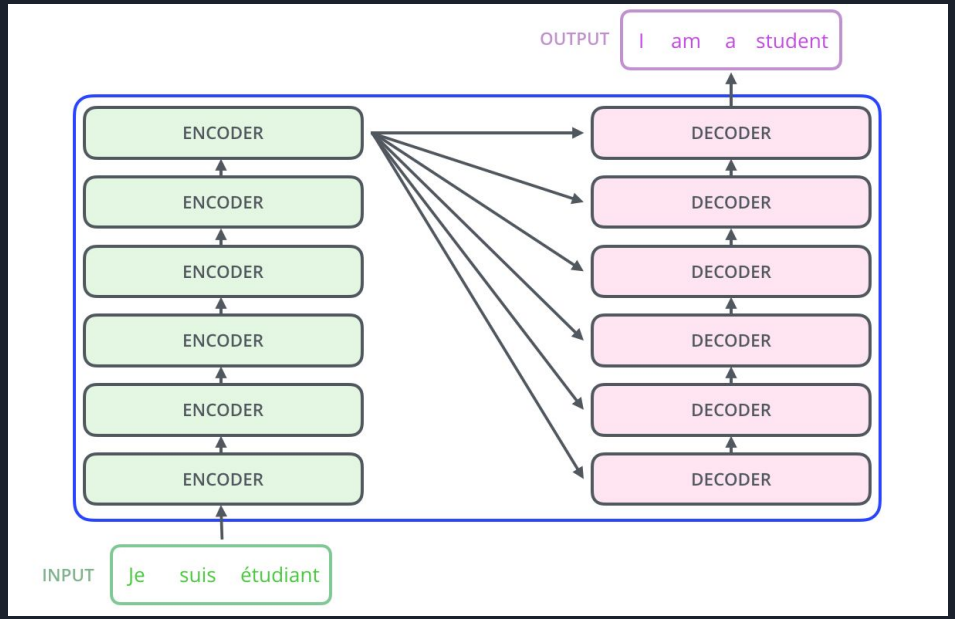

http://jalammar.github.io/illustrated-transformer/

Inside Encoder and Decoder Units

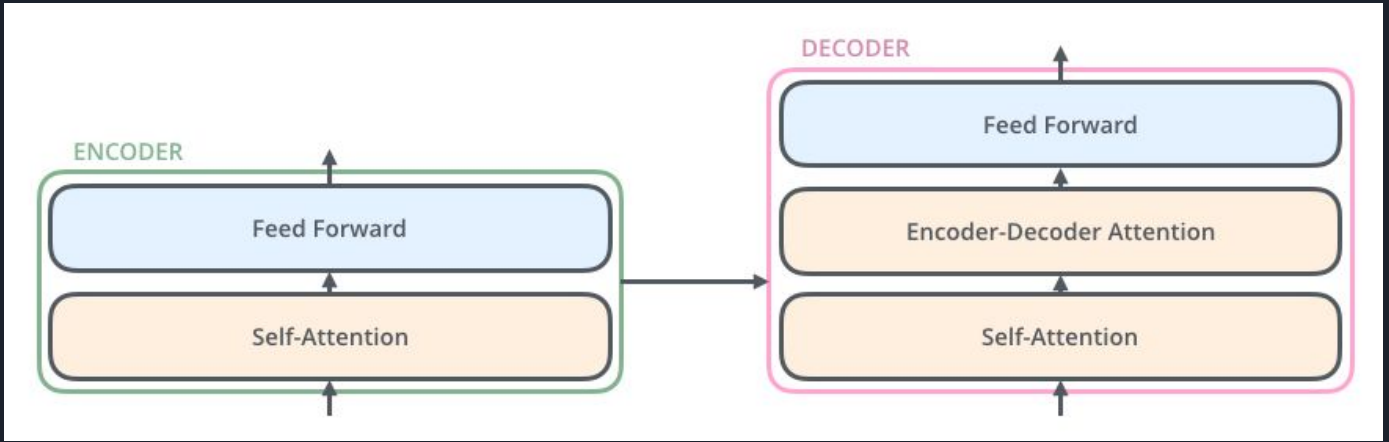

http://jalammar.github.io/illustrated-transformer/

Encoding

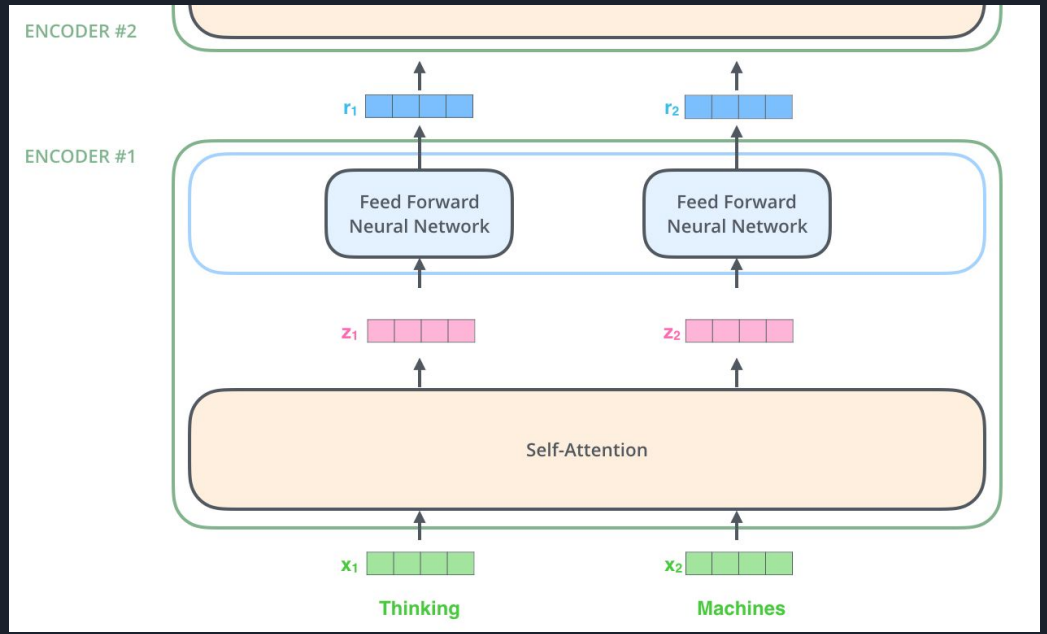

http://jalammar.github.io/illustrated-transformer/

Self-attention

”The animal didn't cross the street because it was too tired”

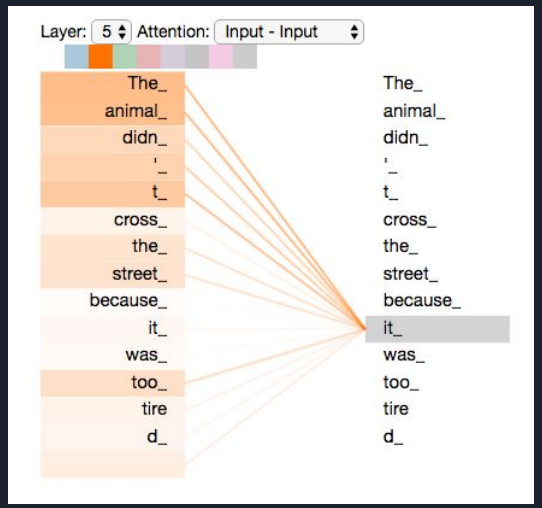

http://jalammar.github.io/illustrated-transformer/

Tensor2Tensor by Google Brain team

Tensor2Tensor, or T2T for short, is a library of deep learning models and datasets designed to make deep learning more accessible and accelerate ML research.

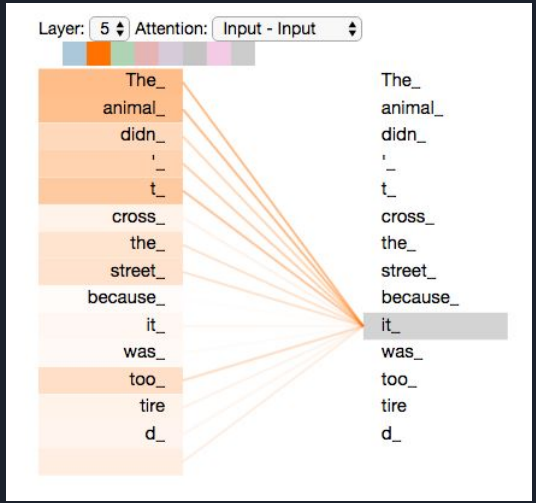

https://colab.research.google.com/github/tensorflow/tensor2tensor/blob/master/tensor2tensor/notebooks/hello_t2t.ipynb

Self-attention in more detail

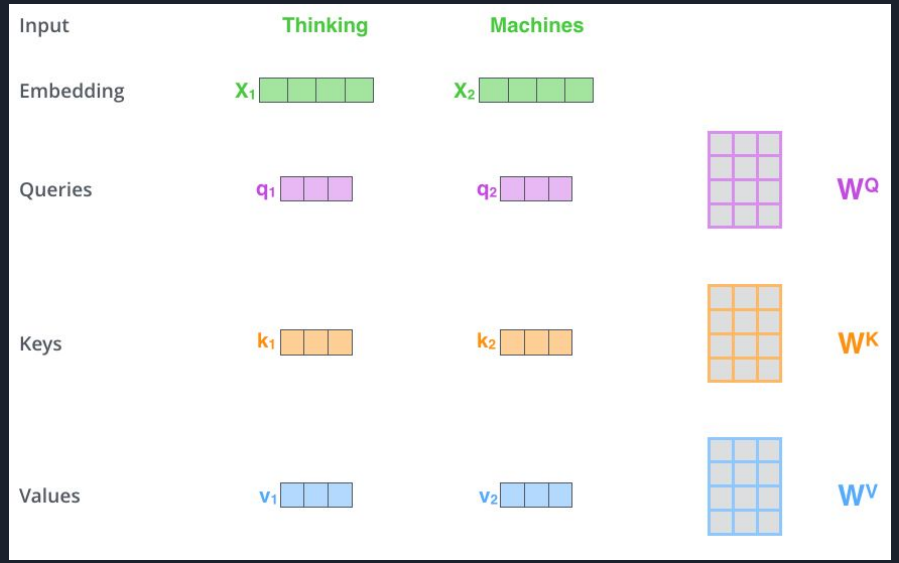

http://jalammar.github.io/illustrated-transformer/

Second step : Score

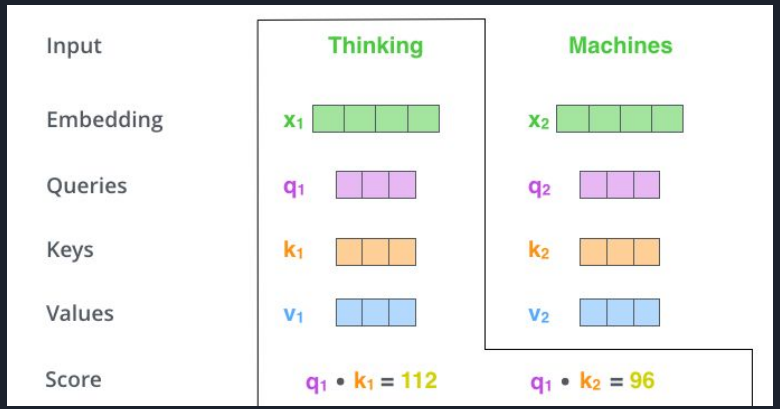

http://jalammar.github.io/illustrated-transformer/

Second step : Score

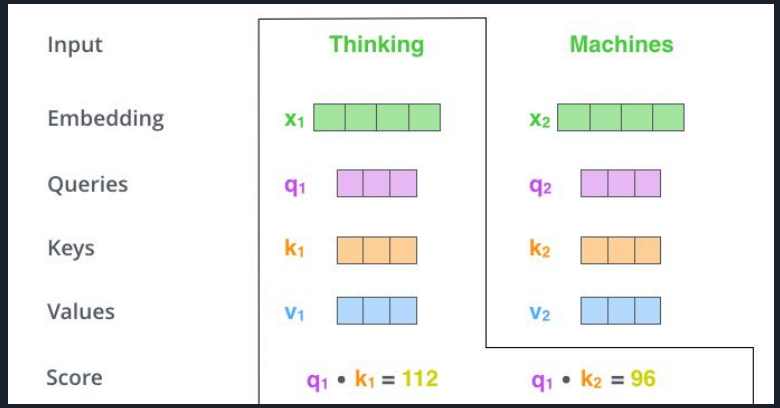

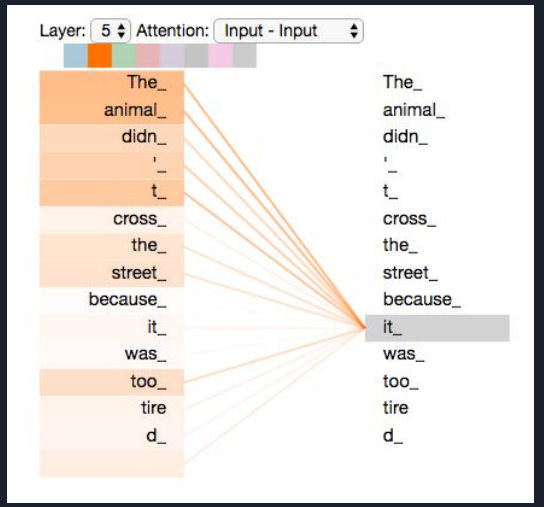

http://jalammar.github.io/illustrated-transformer/

Third step : Divide by 8

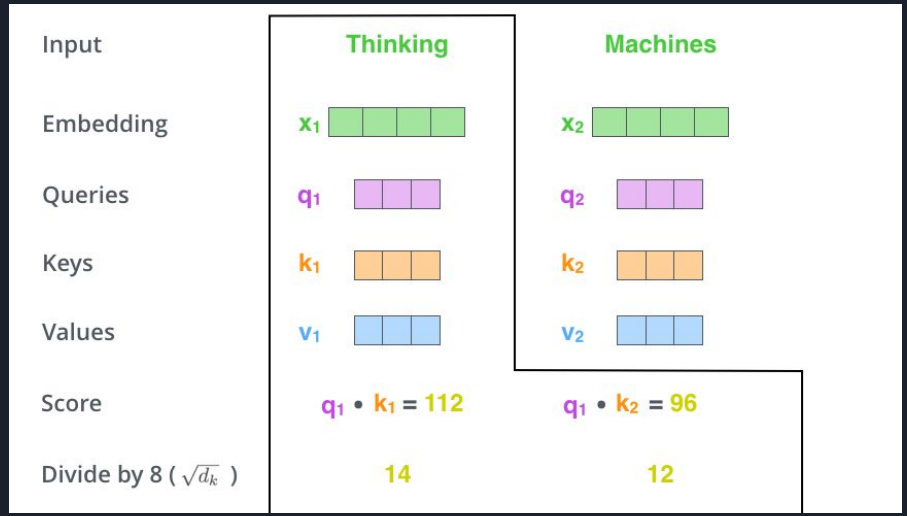

http://jalammar.github.io/illustrated-transformer/

Fourth step : Softmax

Formula >

$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^{K} e^{z_j}}$$

σ = softmax

$\vec{z}$ = input vector

e^{z_i} = standard exponential function for input vector

K = number of classes in the multi-class classifier

e^{z_j} = standard exponential function for output vector

e^{z_j} = standard exponential function for output vector

google.com

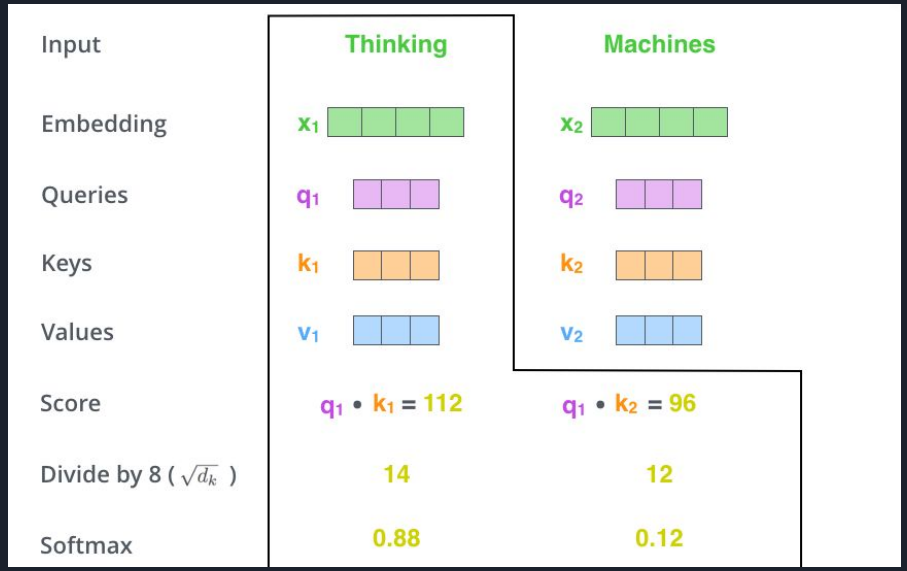

http://jalammar.github.io/illustrated-transformer/

Matrix Representations

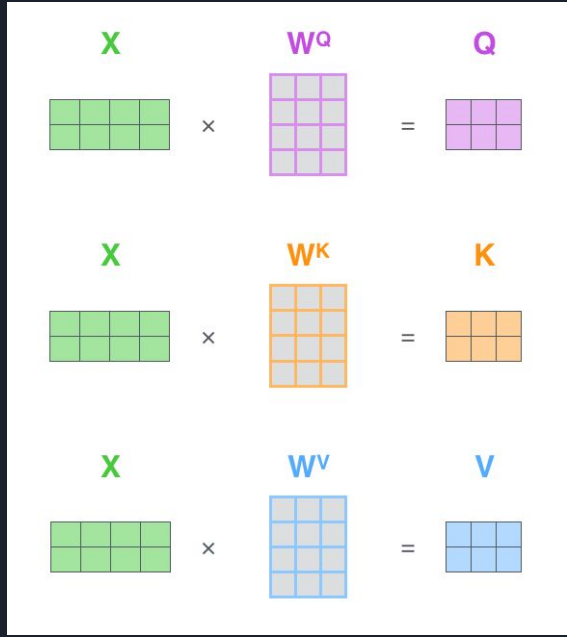

http://jalammar.github.io/illustrated-transformer/

Self-attention Formula

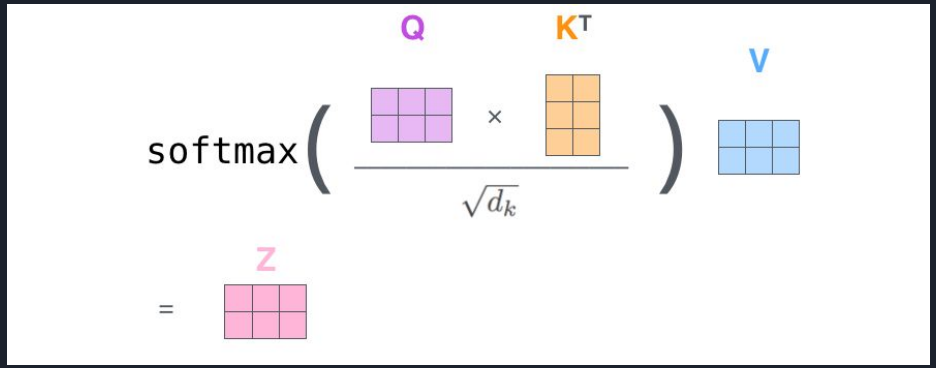

http://jalammar.github.io/illustrated-transformer/

Multi-headed attention

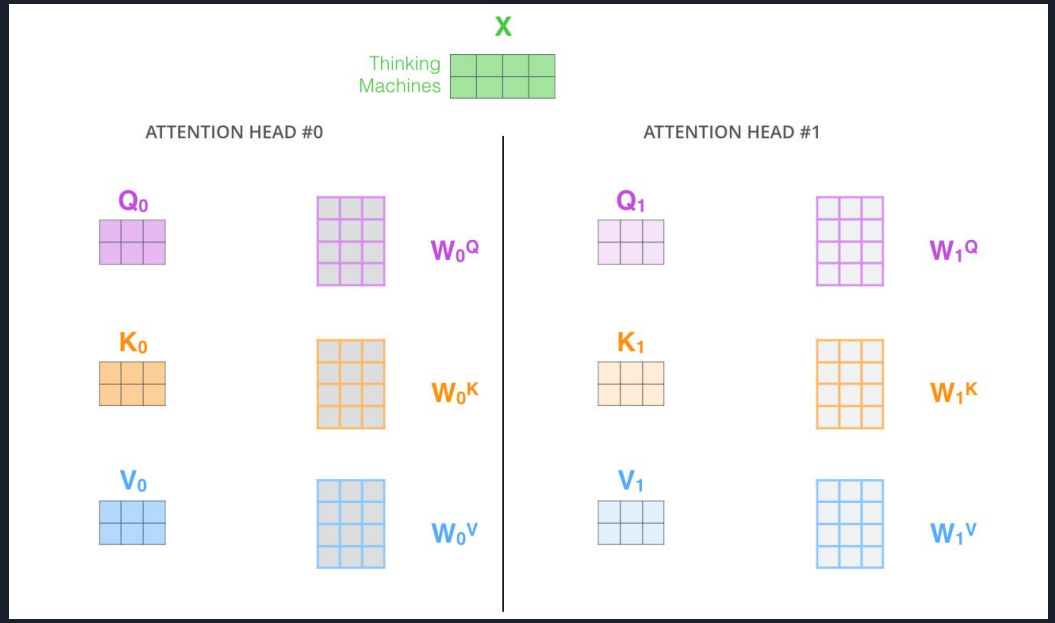

http://jalammar.github.io/illustrated-transformer/

Multi-headed attention

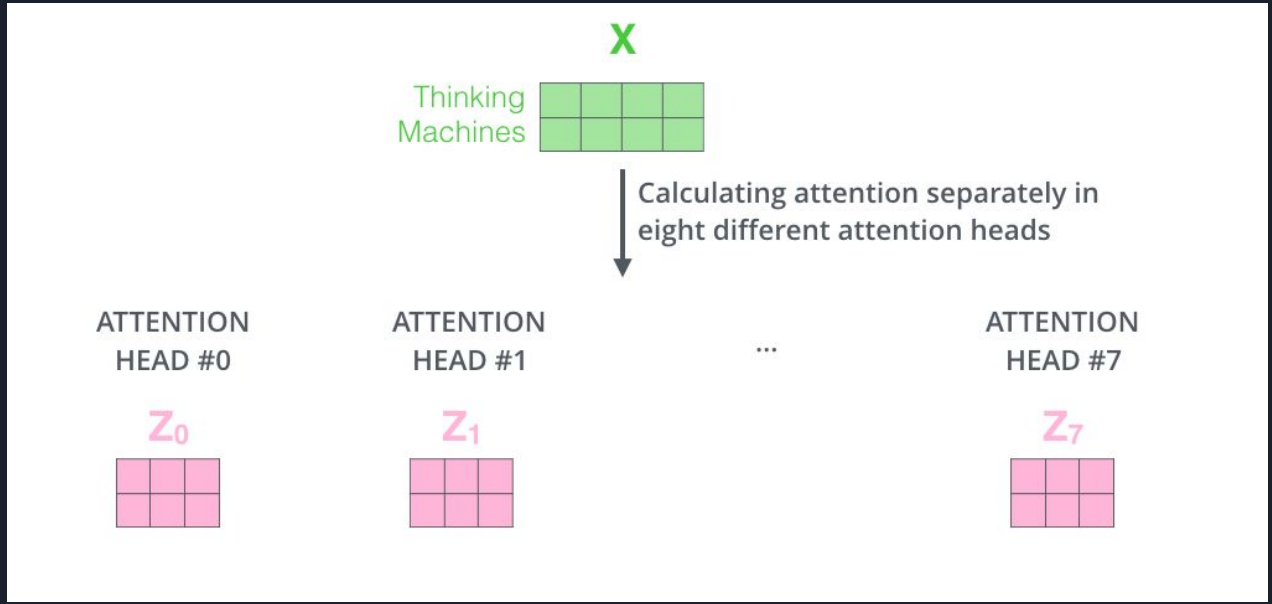

http://jalammar.github.io/illustrated-transformer/

We only need one Z

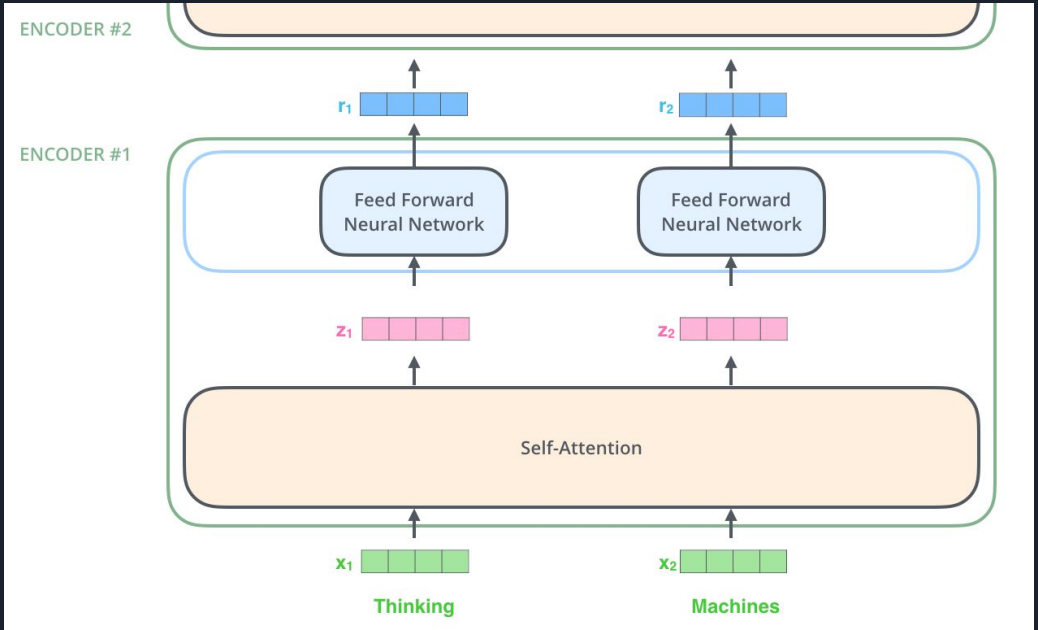

http://jalammar.github.io/illustrated-transformer/

Concatenation and Wo weight matrix

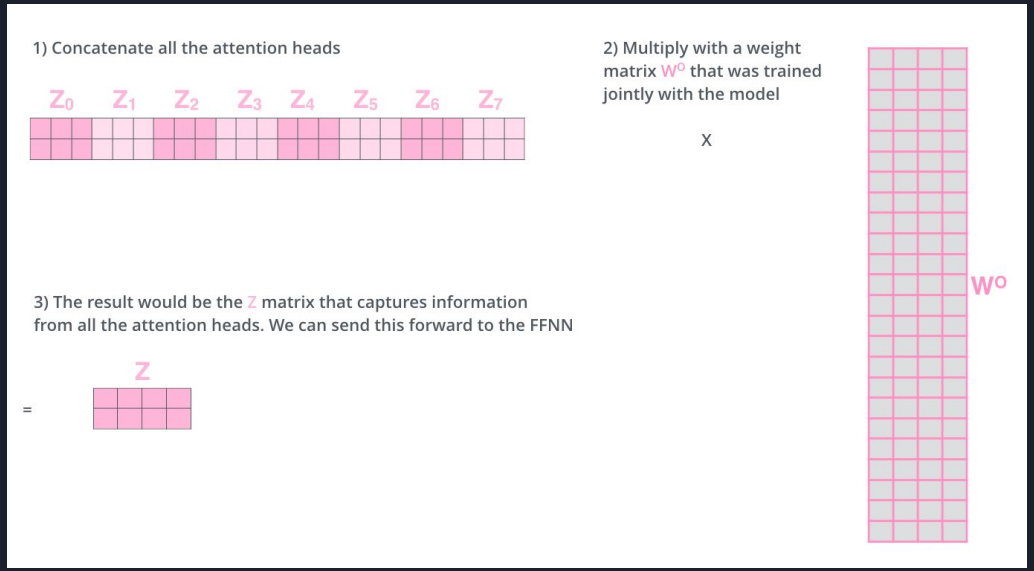

http://jalammar.github.io/illustrated-transformer/

Multi-headed attention visualization

1 Head

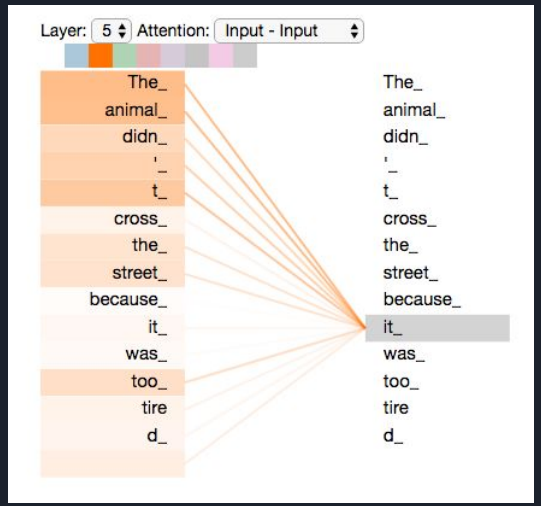

2 Heads

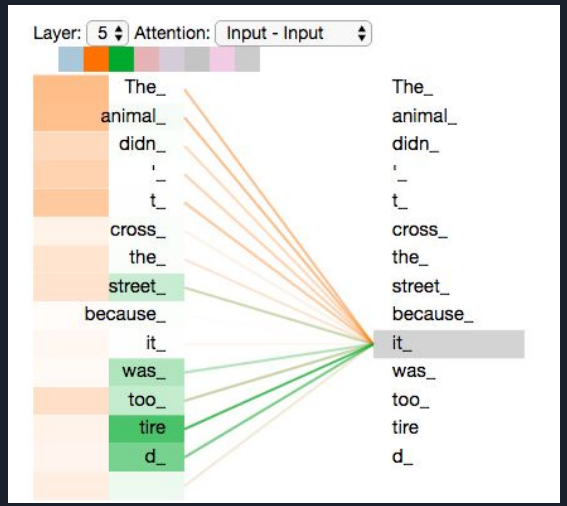

8 Heads

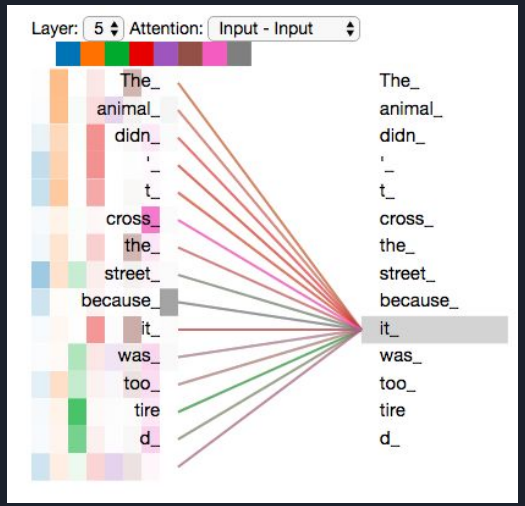

http://jalammar.github.io/illustrated-transformer/

Positional Encoding

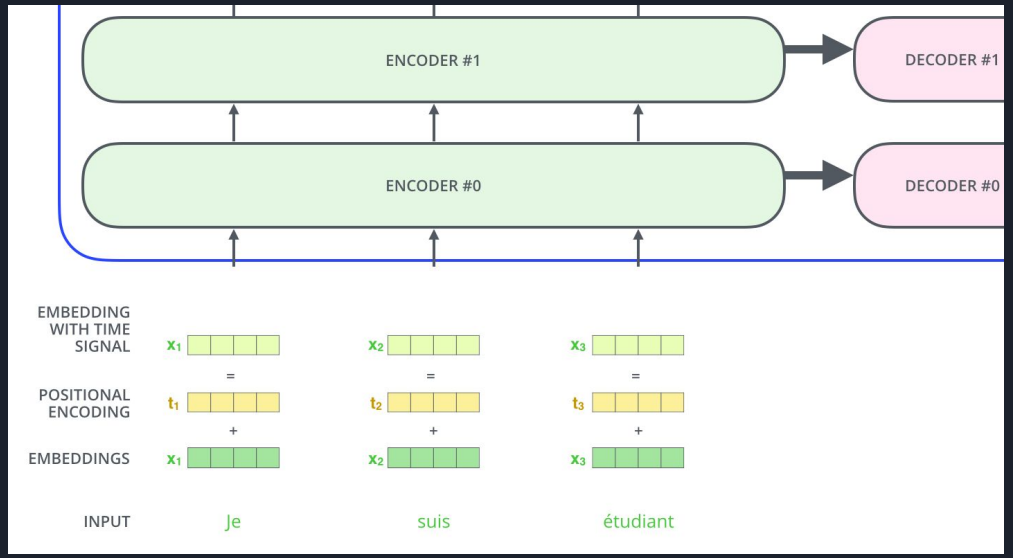

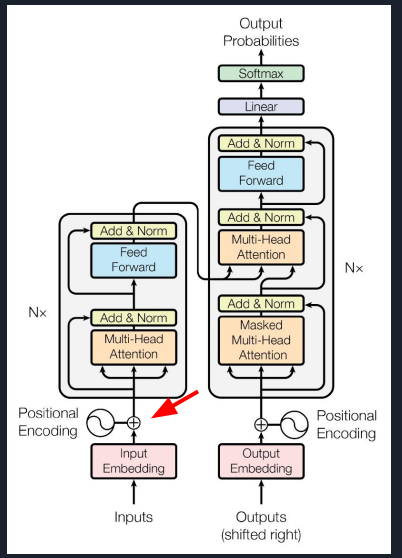

http://jalammar.github.io/illustrated-transformer/

https://arxiv.org/abs/1706.03762

Residual Connections

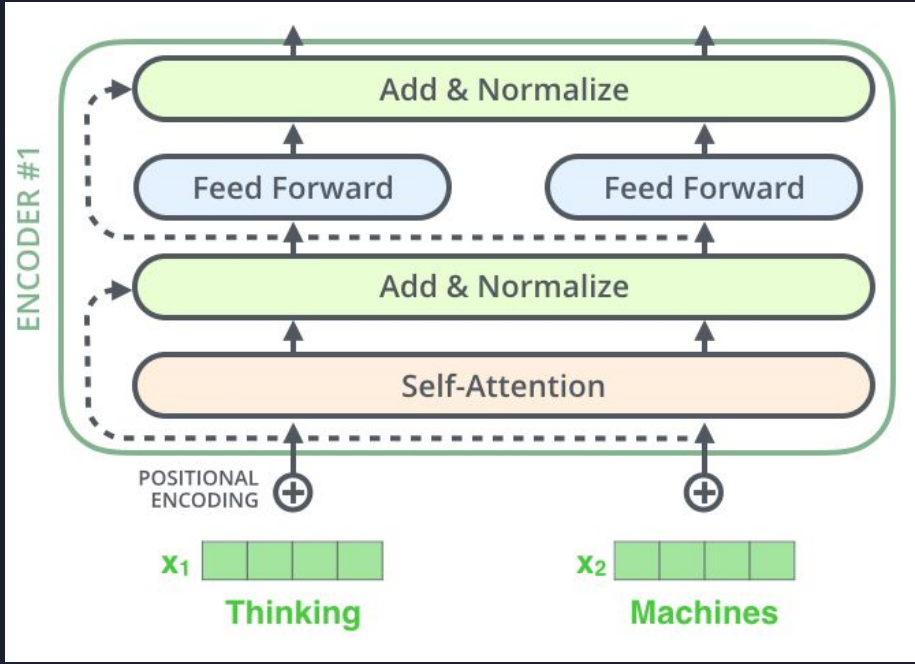

http://jalammar.github.io/illustrated-transformer/

Encoder Decoder Attention

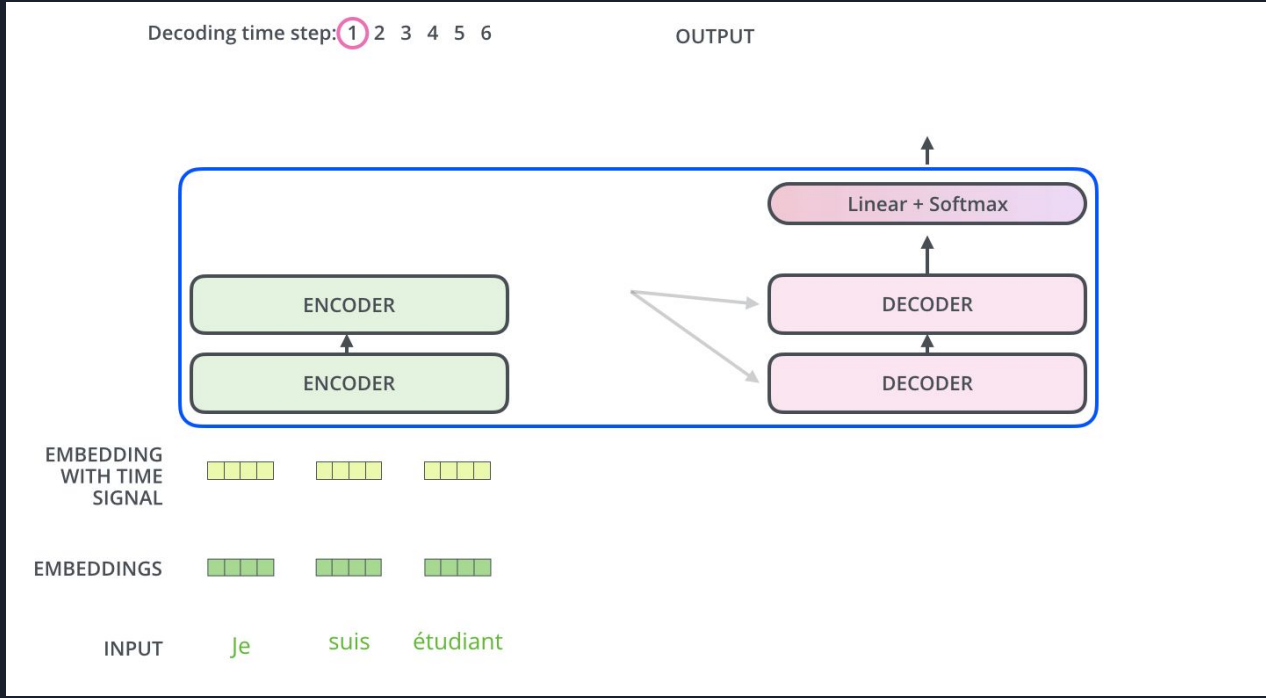

http://jalammar.github.io/illustrated-transformer/

Encoder Decoder Attention

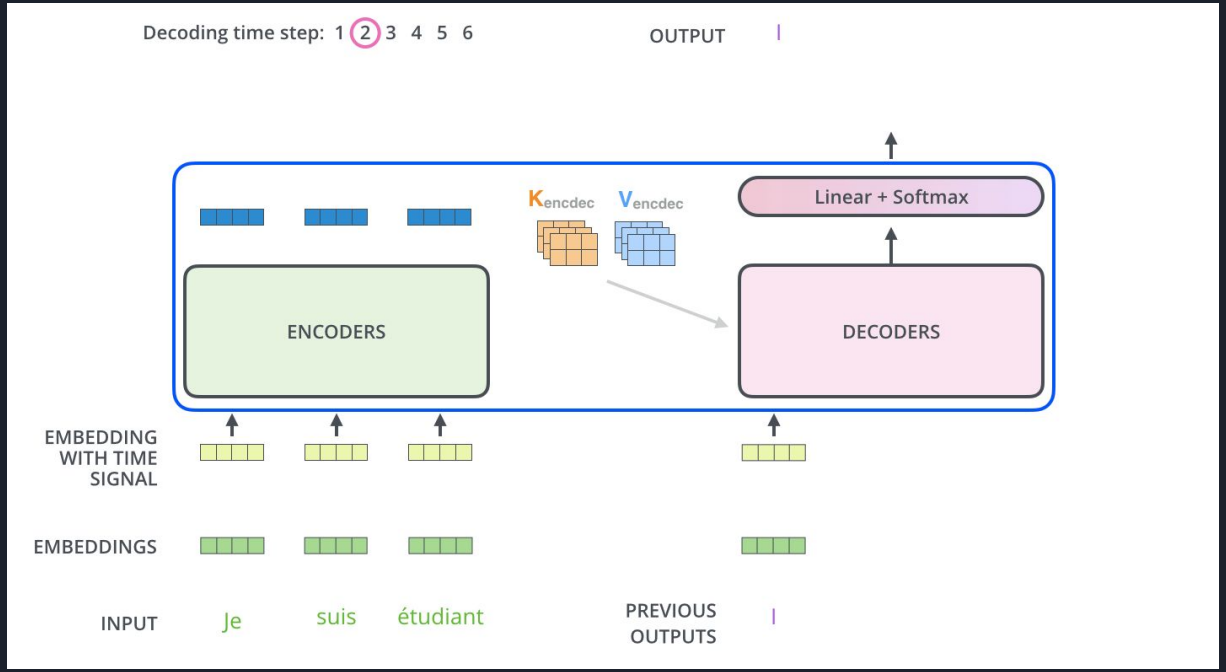

http://jalammar.github.io/illustrated-transformer/

Encoder Decoder Attention

Creates Queries matrix from the layer below it, and takes the Keys and Values matrix from the encoder stack.

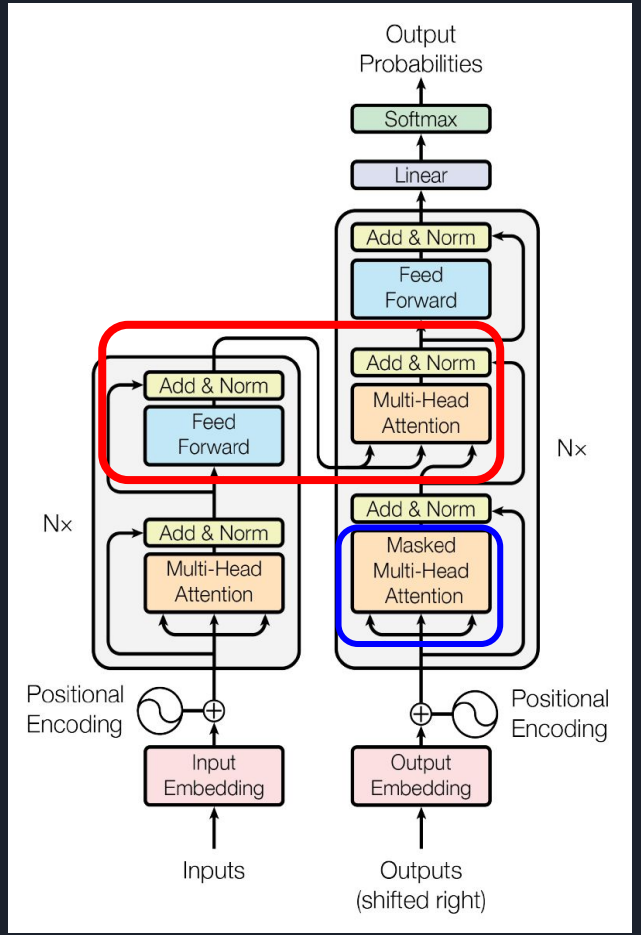

https://arxiv.org/abs/1706.03762

Encoder Attention
Vs
Decoder Attention

Masking

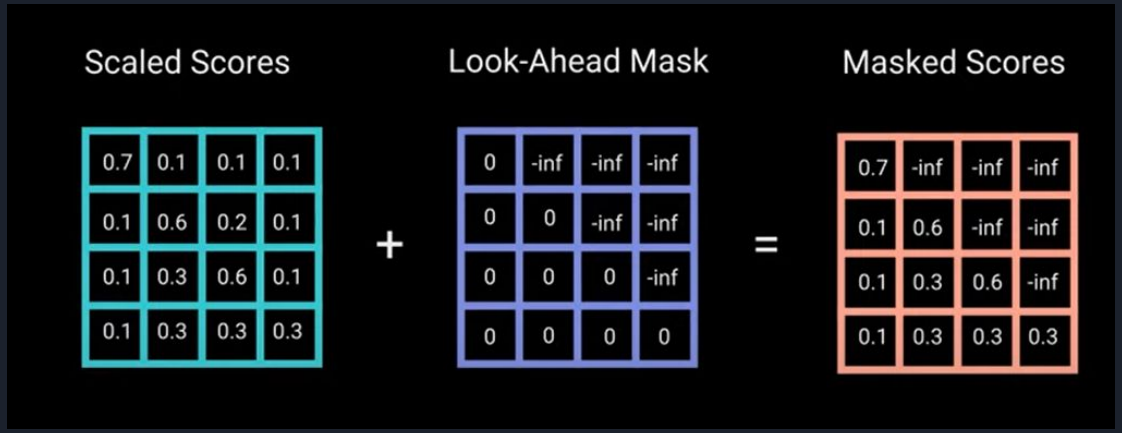

YT - The A.I. Hacker - Michael Phi

The Final Linear and Softmax Layer

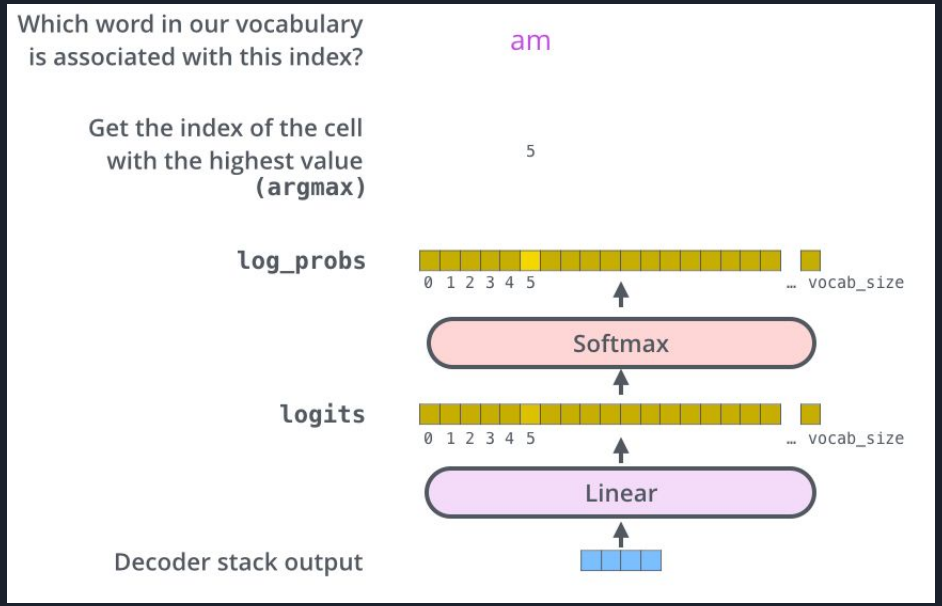

http://jalammar.github.io/illustrated-transformer/

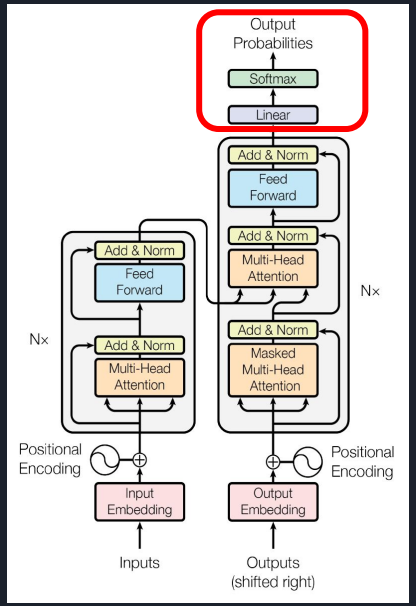

https://arxiv.org/abs/1706.03762

Training

- One Hot Encoding
- Loss Function
- Another Training Method

One Hot Encoding

Output Vocabulary

WORD	a	am	I	thanks	student	<eos>
INDEX	0	1	2	3	4	5

One-hot encoding of the word “am”

0.0	1.0	0.0	0.0	0.0	0.0

http://jalammar.github.io/illustrated-transformer/

Loss Function

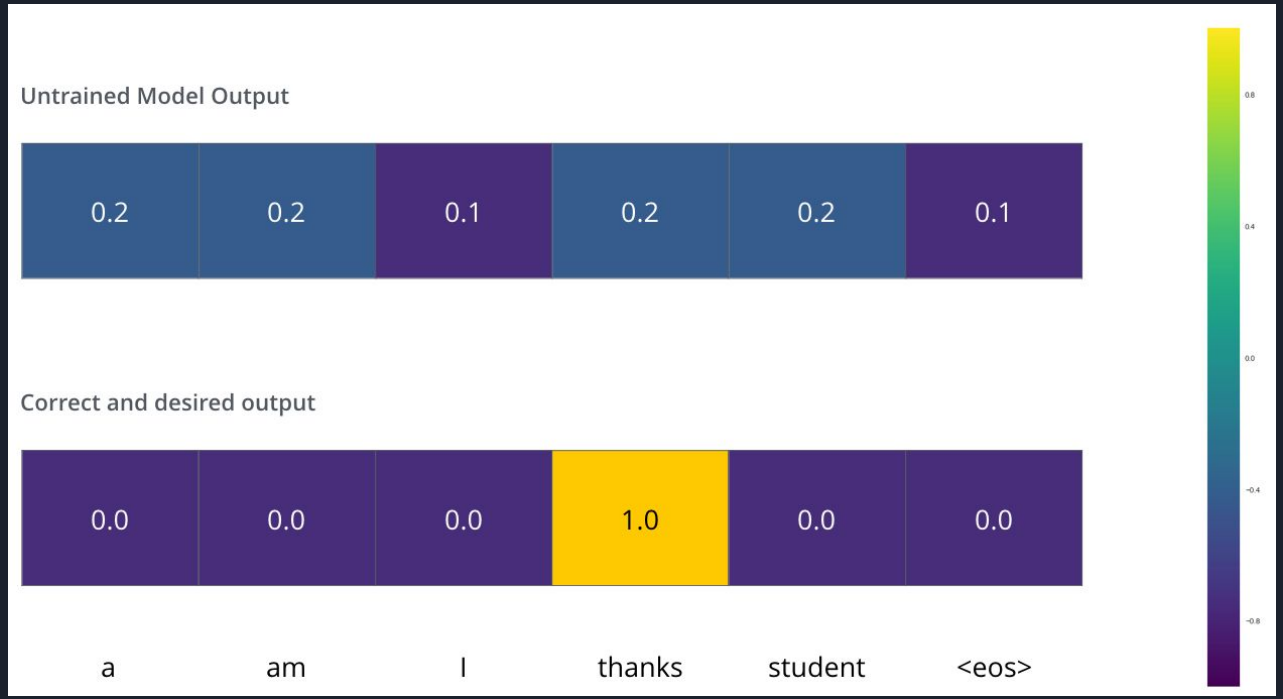

http://jalammar.github.io/illustrated-transformer/

Loss Function

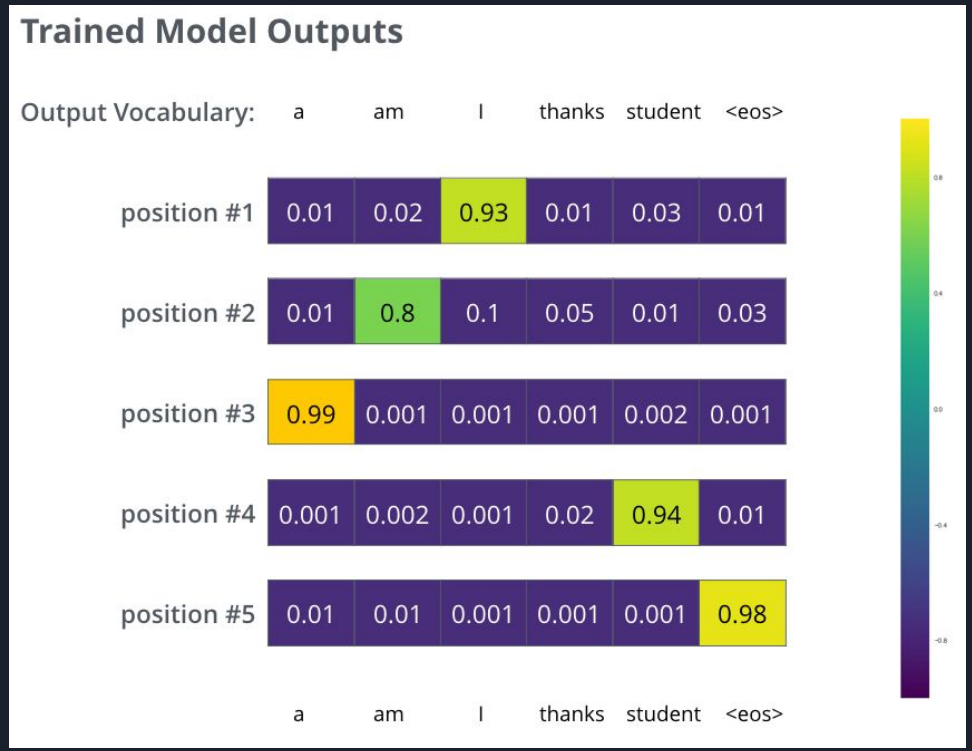

http://jalammar.github.io/illustrated-transformer/

Another Training Method

Another way to do it would be to hold on to, say, the top two words (say, ‘I’ and ‘a’ for example), then in the next step, run the model twice

http://jalammar.github.io/illustrated-transformer/

Resources

Illustrated-transformer

Illustrated Guide to Transformers Neural Network: A step by step explanation

A Deep Dive Into the Transformer Architecture – The Development of Transformer Models

Attention Is All You Need

Transformers and Language Models - YouTube Playlist

Thanks

hasanmahani08@gmail.com

rezadar1378@gmail.com